

TBot-SA₁: 2D-3D Latent World Action Modeling for Generalizable Robot Control

Jiahao Jiang Jianing Zhang Zhenhan Yin Ruidong Chen Sen Wang Zhaoshu Yu
Pengpeng Zeng Xiaofeng Cao Xuanhan Wang[†] Jingkuan Song[†] Heng Tao Shen

School of Computer Science and Technology, Tongji University; Shanghai Innovation Institution

Project Page: <https://magicbot-wvga.github.io>

[†]Project Leader [◇]Corresponding Author




Figure 1: **TBot-SA₁**: A *World-Spatial-Action* (WSA) model for generalizable robot control. Built on 3D world-action bidirectional causal modeling, it is pretrained with only 6K hours of demonstration data, yet achieves strong manipulation performance across simulation and real-robot benchmarks.

Abstract

Embodied foundation models have evolved into two typical paradigms: Vision-Language-Action (VLA) and World-Action-Model (WAM). Although impressive, both paradigms are inherently 2D-centric, lacking innate understanding of 3D physical world dynamics and action-environment causal relationships. This creates a critical discrepancy between 2D visual understanding and 3D-centric embodied interaction, severely compromising the reliability and generalization of existing models in real-world robot control tasks. To address this, we present **TBot-SA₁**, a novel embodied foundation model built on proposed World-Spatial-Action (WSA) paradigm, which learns 2D-3D cross-space transfer from 2D visual observations to 3D robotic behaviors, via the synergistic integration of instruction-aligned visual planning, causal 3D world modeling, and closed-loop 3D action generation. Notably, TBot-SA₁ achieves highly data-efficient pre-training with 6k hours of expert demonstration data (only 1k hours from real robot), while delivering competitive manipulation performance (92% success rate) on RobotWin2.0 simulation benchmark and achieving +20% average boosted performance on real-world robot control tasks compared to state-of-the-art baselines. These results suggest that modeling 3D world-action bidirectional causality in a unified latent space, enables the model to inherently acquire physical commonsense and causal mechanisms of the 3D world, paving a new way for generalizable robot control.

Contents

1	Introduction	3
2	Related Work	4
3	Methodology	5
3.1	Problem Definition and Notation	5
3.2	Model Architecture	6
3.3	Learning Objective	7
3.4	Pre-training Data	8
4	Experiments	9
4.1	Experimental Setup	9
4.2	Main Results on Real-World Tasks	9
4.3	Main Results on Simulation Benchmark	11
4.4	Ablation Study	12
4.5	Qualitative Analysis	13
5	Conclusion	13

1 Introduction

Creating versatile robots to match or exceed human-level physical intelligence has long been a core pursuit in embodied AI. Serving as the cornerstone of human physical intelligence, the innate capacity for spatial world modeling enables humans to integrate 2D visual understanding and task objectives, predict physical dynamics in 3D space, and execute precise motor actions [12, 14]. Such 2D-to-3D spatial reasoning capability is equally vital for versatile robots, as physical interactions in the real world inherently occur in 3D space and follow deterministic physical laws. Accordingly, it is imperative to develop an embodied foundation model that can master intrinsic causal relationships between spatial environment evolution and robot behaviors, as well as physical common sense behind interactions in the 3D world.

In recent years, multi-modal large models [11, 3, 43, 34] have achieved remarkable progress and, meanwhile, demonstrated a promising technical roadmap toward beyond human-level intelligence through training on web-scale data. Following this direction, modern embodied foundation models have gradually evolved into two typical paradigms: Vision-Language-Action (VLA) models [22, 49, 35, 30, 7, 6, 16, 47] and World Action Models (WAM) [4, 26, 9, 17, 42, 39, 15]. VLA models construct semantic-centric robotic policies by augmenting pre-trained vision-language models (VLMs) with a lightweight action expert. This design allows VLAs to inherit powerful open-world visual understanding capabilities from pre-trained VLMs, while the appended action expert generates robot actions conditioned on semantic-level visual comprehension, following an "understand-then-execute" paradigm. In contrast, WAM learns imagination-centric robot policies based on a predictive video generation model that acts as a vision planner. Based on this "imagine-then-execute" paradigm, its action expert functions as an inverse dynamics model (IDM), which generates robotic control signals corresponding to the planned visual frames.

Despite the impressive progress of current VLAs and WAMs, two fundamental challenges still impede their generalization to real-world scenarios. **First**, generative vision-language learning in VLAs and predictive video learning in WAMs are inherently 2D-centric [27, 46, 28]. Both paradigms operate primarily on 2D visual signals, failing to explicitly model the 3D physical dynamics of the environment and the bidirectional causal relationships between robot actions and scene changes. This renders them unreliable in real-world tasks that require 3D spatial reasoning and robust adaptation to environmental changes. **Second**, learning generalizable embodied foundation models heavily depends on real-world expert demonstration data, which is costly and time-consuming to acquire at scale [5]. This data scarcity further prevents models from learning the generalizable physical commonsense behind demonstrations, which is essential for developing versatile robots that can adapt to unseen scenarios [5, 4]. These challenges motivate us to study a significant question:

How to learn unified 2D–3D world–action bidirectional causality in a shared latent space for constructing data-efficient and generalizable embodied foundation models under constrain of limited expert demonstration data?

In this work, we present **TBot-SA₁**, a novel embodied foundation model built on data-efficient World-Spatial-Action (WSA) paradigm. Specifically, the WSA breaks the inherent 2D-centric limitations of prevailing paradigms via integration of three essential capabilities. First, **at the semantic understanding level**, it enables instruction-aligned 2D visual planning via subgoal image prediction, empowering a robot to accurately parse the underlying intent of task instructions. Second, **at the causal reasoning level**, it incorporates a causal 3D scene dynamics mechanism that explicitly models the evolution of 3D scenes conditioned on robot actions, endowing a robot with the ability to capture the intrinsic causal relationships between environmental changes and robot behaviors. Third, **at the closed-loop execution level**, it learns a 3D inverse dynamics model that maps the planned 3D scene states to executable robot control signals, fully closing the loop between 3D spatial world modeling and real-world physical interaction. To this end, the TBot-SA₁ is implemented via a Mixture-of-Transformers (MoT) architecture with 3D-centric bidirectional causal attention, enabling efficient joint learning of all three capabilities. In summary, our main contributions are three-fold:

- **Robot Learning Paradigm:** We propose World-Spatial-Action (WSA) modeling, a novel robot learning paradigm that unifies instruction-aligned semantic understanding, causal 3D world modeling, and closed-loop physical execution into a single model, addressing the inherent 2D-centric limitations of prevailing VLA and WAM paradigms.




Figure 2: **Representative Model Paradigms:** (a) 2D-centric VLA models learn unidirectional causality from visual semantics to robot action; (b) 2D-centric World Action Model learns unidirectional causality from 2D visual dynamics to robot actions; (c) 3D-centric world action model learns unidirectional causality from 3D dynamics to robot action; (d) World-Spatial-Action model unifies bidirectional causality between 3D dynamics and robot action into a single world.

- **Embodied Foundation Model:** We present TBot-SA₁, a data-efficient embodied foundation model implementing the WSA paradigm via a Mixture-of-Transformers (MoT) architecture, enabling bidirectional causal reasoning between 3D world dynamics and robot actions in a shared latent space.
- **Data-efficient Learning and Strong Performance:** We pre-train TBot-SA₁ with 6k hours of demonstration data (only 1k hours from real robot), and it delivers competitive manipulation performance (e.g., 92% success rate on RobotWin2.0) on simulation benchmark and achieves +20% boosted performance on real-world robot control tasks.

2 Related Work

The core goal of embodied foundation models (EFMs) is to establish a robust, stable multi-modal mapping from robot observations and task instructions to executable robot actions [44, 41]. As illustrated in Fig. 2, mainstream research on EFMs has converged into two dominant paradigms: Vision-Language-Action models (VLAs) and World-Action-Models (WAMs). In the following sections, we provide a systematic literature review of these two research directions.

Vision-Language-Action: As the most widely adopted paradigm for robot manipulation, VLA models [35, 7, 6, 37, 47] are generally built on pretrained vision-language models. By inheriting rich semantic priors from VLMs [11, 3], they successfully transfer open-world visual-semantic knowledge to the physical robot domain via end-to-end fine-tuning on robot demonstration data. Early representative works such as RT-2 [49] and OpenVLA [16] adopt an autoregressive strategy, which discretizes continuous robot actions and formulates policy learning as a sequence modeling task. To address the limitation of discrete tokenization in fine-grained continuous control, recent state-of-the-art works [7, 6, 5] have shifted to generative frameworks based on conditional flow matching, which directly model the continuous-valued action, significantly improving the precision and stability of robot control. In addition, several works [48, 28] augment VLA with 3D geometry priors by incorporating extra 3D information such as depth or 3D gaussian splatting, demonstrating the effectiveness of 3D representation in robot learning. However, mainstream VLA policies are essentially reactive frameworks that make decisions based only on current observations, lacking

predictive planning and 3D causal reasoning capabilities, which makes them unreliable in real-world manipulation tasks that require 3D dynamics modeling.

World-Action-Model: In a different line from VLAs, WAMs follow the chain-of-thought philosophy and formulate robot policies as an imagine-then-execute procedure. Guided by this principle, numerous works [26, 39, 15, 42, 23] start with video models, inheriting rich commonsense knowledge from large-scale video generation pre-training. They then use these video models as world dynamics simulators, enhancing the reliability of robot foundation models via co-training of future frame prediction and action generation. For example, Motus [4] demonstrates that explicitly modeling future frame prediction significantly improves the generalization of robot manipulation, while Fast-WAM [42] proves that improved manipulative generalization comes from the co-learning rather than explicit video prediction. Furthermore, recent efforts [27, 45, 19] have further extended WAM paradigm to 3D space by modeling future geometry prediction. For instance, GeoPredict [27] leverages predictive kinematics and 3D Gaussian geometry to improve the manipulation performance, and PointWorld [14] scales 3D world models for in-the-wild robotic manipulation. However, existing 3D WAMs still follow a unidirectional "predict-then-execute" pipeline, failing to unify 3D world dynamics modeling and inverse dynamics into a single latent space for bidirectional causal learning. They cannot explicitly model how robot actions affect 3D space evolution and, in turn, how 3D space changes should guide action generation, which is the core focus of our work.

Compared with prior studies, our work is most closely related to WAMs but takes a further step: instead of directly co-learning video generation and robot control, our method unifies action-conditioned 3D prediction and 3D inverse dynamics into a single model. This mechanism enables our model to exploit the intrinsic causal consistency: it not only learns how 3D robot actions effect the evolution of the 3D world but also determines the appropriate action plan given variations in the 3D world.

3 Methodology

In this section, we first formalize the robot control problem in Section 3.1. We then introduce the overall architecture of TBot-SA₁ in Section 3.2, followed by detailed learning objectives of the World-Spatial-Action (WSA) modeling mechanism in Section 3.3. Finally, we present the pre-train data recipe in Section 3.4.

3.1 Problem Definition and Notation

In the context of robot control, the goal of embodied foundation model is to learn a robot policy that predicts next robot actions from current observations. Formally, at each timestep t , the observations include task instruction termed l , visual observation denoted as v_t , and robot proprioception termed s_t . In general, the robot action is a H -step chunk termed $A_t := (\mathbf{a}_{t+1}, \dots, \mathbf{a}_{t+H})$, a short-term action trajectory with a sequence of end-effector pose or joints of a robot in next H timesteps. To train a robot policy π_θ parameterized by θ on diverse robot demonstrations $\mathcal{D}_{\text{dem}}$, it is required to maximize the the standard likelihood objective:

$$\max_{\theta} \mathbb{E}_{(v_t, s_t, A_t, l) \sim \mathcal{D}_{\text{dem}}} \log \pi_\theta(A_t \mid v_t, s_t, l). \quad (1)$$

As illustrated in Fig. 2, previous works formulate robot control problem as vision-language-action or world-action policy, which can be regarded as a joint distribution derived from the standard objective.

2D Vision-Language-Action Policy. The 2D-centric VLA aims to directly predict executable actions conditioned on the current 2D observation $\mathbf{O}_t := \{v_t, s_t, l\}$. It learns 2D semantic-action unidirectional causality, which is formulated as a posterior distribution:

$$\mathbf{A}_t \sim p(a_{t+1:t+H} \mid v_t, s_t, l). \quad (2)$$

2D World-Action-Model Policy. The 2D-centric WAM couples predictive visual generation with next action generation, which jointly predicts executable actions and N subgoal images $\mathbf{V}_t := (\mathbf{v}_{t+1}, \dots, \mathbf{v}_{t+N})$ from the current observation $\mathbf{O}_t$. It learns 2D world-action unidirectional causality, which can be formulated as a joint prediction distribution:

$$\mathbf{A}_t \sim p(a_{t+1:t+H}, v_{t+1:t+N} \mid v_t, s_t, l). \quad (3)$$




Figure 3: Overview of TBot-SA₁. It adopts a Mixture-of-Transformers structure with three complementary experts: a 2D Spatial Expert for instruction-aligned visual planning, a 3D Spatial Expert for action-conditioned 3D world prediction, and a 3D Action Expert for 3D robot action generation. A bidirectional causal attention mechanism enforces dependency rules to unify 2D-3D visual consistency and 3D bidirectional action–scene causality within a shared latent space, enabling joint learning of visual planning, 3D world modeling, and 3D inverse dynamics model for robot control.

3D World-Action-Model Policy. Instead of modeling 2D world-action joint distribution, the 3D-centric WAM focuses on the unidirectional causality between evolution of scenes and robot action in 3D space. Specifically, given 2D observation $\mathbf{O}_t$ and 3D observation g_t , it learns to predict K subgoal 3D scenes $\mathbf{G}_t := (g_{t+1}, \dots, g_{t+K})$ for robot control. Formally, this models a joint prediction distribution as followed:

$$\mathbf{A}_t \sim p(a_{t+1:t+H}, g_{t+1:t+K} \mid g_t, v_t, s_t, l). \quad (4)$$

3D World Spatial Action Policy. Different from previous works, we formulate robot control as 3D-centric world-spatial-action learning, which jointly models 2D semantic-world unidirectional causality and 3D world-action bidirectional causality. Given the current observations $\mathbf{O}_t$, the WSA models a joint distribution upon three prediction tasks that are relevant to robot control, including 3D-conditioned subgoal image prediction, action-conditioned 3D scene prediction and 3D-conditioned action generation:

$$\mathbf{A}_t \sim p(\mathbf{V}_t, \mathbf{G}_t, \mathbf{A}_t \mid \mathbf{O}_t) = p(\mathbf{V}_t \mid \mathbf{O}_t, \mathbf{G}_t) \cdot p(\mathbf{G}_t \mid \mathbf{O}_t, \mathbf{A}_t) \cdot p(\mathbf{A}_t \mid \mathbf{O}_t, \mathbf{G}_t). \quad (5)$$

3.2 Model Architecture

Inspired by the recent works [4, 6], TBot-SA₁ is built upon a unified multi-modal transformer, which leverages mixture-of-transformer (MoT) architecture to enable seamless integration of 2D visual generation, future 3D prediction and robot action generation. As illustrated in Fig. 3, it involves three major components: 2D spatial expert, 3D spatial expert and 3D action expert.

2D Spatial Expert (2D-SE): Given a human instruction, generating a sequence of visual changes in 2D space enhances understanding of human intent. The 2D-SE $f_{2D}(\cdot)$ serves as the instruction-

aligned semantic understanding branch of TBot-SA₁. It is initialized from a pre-trained vision-language model such as QWen3-VL [2] or Wan2.1 [31], thereby inheriting strong vision-language semantic priors. The goal of 2D-SE is to predict future dynamics that describe how 2D visual appearance changes so that follow the instruction. Specifically, it predicts highly abstracted tokens $h_v = f_{2D}(O_t, G_t)$ for subgoal images, instead of RGB pixels. This design enhance visual understanding capability of the 2D-SE, as it is required to focus on essential features of an image rather than redundant pixel information.

3D Spatial Expert (3D-SE): Given a sequence of executable robot actions, predicting evolution of world dynamics in 3D space reflects causal effect of the robot behavior. Under this context, the 3D-SE $f_{3D}(\cdot)$ conducts a internal mechanism for 3D causal world modeling. To capture essential feature of 3D scenes, the 3D-SE predicts latent representations $h_g = f_{3D}(O_t, A_t)$ for subgoal 3D scenes. We construct it with a transformer-based module of the same depth as 2D-SE, thereby the architecture of 3D-SE can be directly derived from that of 2D-SE.

3D Action Expert (3D-AE): Given a demonstration of 3D space evolution, predicting next robot action tells why and how robot to behave in 3D space. The 3D-AE $f_{act}(\cdot)$ aims to predict a full action chunk conditioned on both current observations and future 3D latent representations. Specifically, we instantiate 3D-AE as denoising diffusion transformer. Given noisy actions $\hat{A}_t = \tau A_t + (1 - \tau)\omega$, where $\omega \sim \mathcal{N}(0, 1)$ and $\tau \in [0, 1]$, it employs iterative cross-attention and denoising procedure to predict latent action representations $h_{act} = f_{act}(O_t, G_t, \hat{A}_t)$. This yields a 3D inverse dynamics formulation, where 3D-AE maps predicted 3D latent evolution back to executable action.

3D-centric World-Action Causal Attention: Three experts are coupled within the mixture-of-transformer architecture via a shared attention mechanism. Formally, let h_t denote tokens of current observation, which are extracted via 2D-SE. As illustrated in Fig. 3, we adopt a 2D-3D causal attention mechanism, which defines the following dependency rules: (1) All predicted tokens (h_v, h_g, h_{act}) can attend to the current observation tokens h_t , ensuring all predictions are grounded in the current state of the robot and the environment. (2) The visual tokens of future subgoal images h_v can attend to the 3D scene tokens h_g , encouraging the model to learn 2D-3D visual consistency and ground semantic visual predictions in physically consistent 3D geometry. (3) The 3D scene tokens h_g and action tokens h_{act} have full bidirectional attention access to each other, enabling joint learning of action-conditioned 3D dynamics prediction and 3D scene-aware action generation. This bidirectional attention is the key to bidirectional causal modeling of the WSA paradigm.

3.3 Learning Objective

TBot-SA₁ is optimized end-to-end with three joint training objectives, which correspond to the three factorized terms in the WSA joint distribution (Equ. 5) and align with the three core capabilities of the WSA paradigm:

3D World-Aware Visual Planning (3D WA-VP). We adopt a MSE loss over the predicted visual tokens, to minimize the discrepancy between the predicted subgoal image tokens and the ground-truth tokens extracted via pre-trained tokenizer. Formally, the loss is defined as:

$$\mathcal{L}_{2D} = \mathbb{E}_{(O_t, G_t, V_t) \sim \mathcal{D}_{dem}} \|f_{2D}(O_t, G_t) - f_{enc}(V_t)\|_2^2, \quad (6)$$

where $f_{enc}(\cdot)$ denotes a pre-trained VAE-based tokenizer (e.g., COSMOS [1]).

Action-Conditioned 3D World Modeling (3D AC-WM). We adopt MSE loss over the predicted 3D scene latent representations, to minimize the error between the predicted 3D scene evolution and the ground-truth 3D geometric states from the demonstration data. The loss is formalized as:

$$\mathcal{L}_{3D} = \mathbb{E}_{(O_t, A_t, G_t) \sim \mathcal{D}_{dem}} \|f_{3D}(O_t, A_t) - f_g(G_t)\|_2^2, \quad (7)$$

where $f_g(\cdot)$ denotes a pre-trained 3D foundation model, e.g., Depth-Anything [20].

3D Inverse Dynamics Modeling (3D IDM). Following previous works [27, 6], we adopt flow matching loss [21] for continuous action generation. This loss minimizes the error between the

Table 1: Data mixture used for pretraining.

Data source	Type	Num. frames	Num. tasks	Sampling weight
InternData-A1 [32]	Sim.	396M	70	0.64
RoboTwin [10]	Sim.	17M	50	0.08
AgiBot-World [8]	Real	206M	217	0.18
RoboChallenge [36]	Real	5M	30	0.02
EgoDex [13]	Human	68M	194	0.08

predicted flow field and the ground-truth one, conditioned on the current observation and the 3D scene state evolution. The loss is defined as:

$$\mathcal{L}_{\text{ACT}} = \mathbb{E}_{(O_t, G_t, A_t) \sim \mathcal{D}_{\text{dem}}} \left\| \omega - A_t - f_{\text{act}}(O_t, G_t, \hat{A}_t) \right\|_2^2. \quad (8)$$

This trains 3D-AE to predict flow field ($\omega - A_t$), thereby facilitating coherent action generation. The overall learning objective is the sum of the three losses:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{2D}} + \mathcal{L}_{\text{3D}} + \mathcal{L}_{\text{ACT}}. \quad (9)$$

We adopt two-stage strategy to optimize the TBot-SA₁: pre-training on diverse data sources to acquire real-world behavior priors in first stage; post-training on specific manipulation tasks to adapt learned priors to target embodiment in second stage. Different from previous works [9, 4], TBot-SA₁ is trained by co-learning of three capabilities in both pre-training and post-training stage, rather than using a hierarchical manner.

3.4 Pre-training Data

High-quality embodied data is the cornerstone of pre-training generalizable embodied foundation models that can acquire versatile manipulation skills for flexible real-world robot control. However, due to the prohibitive cost and safety constraints of human-robot teleoperation, collecting large-scale high-quality real-world robot data at scale remains challenging [5]. To address this limitation, alternative embodied data sources such as simulation data and egocentric human data have emerged as critical complements. This gives rise to a multi-sourced data pyramid [5, 4], where data quantity decreases from the bottom to the top. It suggests that each data layer exhibits distinct, complementary strengths and inherent limitations, and no single data source alone can support the development of generalizable embodied foundation model: human data (at the bottom of the pyramid) provides the "what" and "why" of tasks, simulation robot data (at the middle) offers the "how" for robot control, and real-world robot data (at the top) bridges the "reality gap" between simulation and physical environments [40]. Their synergistic integration is therefore not merely beneficial but essential for pre-training robust, versatile, and human-aligned robot foundation models that can generalize across tasks, platforms, and physical environments. Accordingly, as illustrated in Table 1, we pre-train the model on a mixture of heterogeneous data sources:

- **Simulation data:** We collect simulated synthetic data are from InternData-A1 [32] and RoboTwin2.0 [10], providing diverse and low-cost manipulation trajectories for learning scene dynamics, object interaction, and coordinated control.
- **Human data:** We include egocentric human manipulation videos from EgoDex [13], which provide rich hand-object interaction priors and further enhance the model’s understanding of manipulation behaviors and scene evolution.
- **Real-world data:** The real-world robot data are drawn from RoboChallenge [36] and AgiBot-World [8], which help reduce the sim-to-real gap and improve robustness in physical environments.

The mixture of data sources covers 10 different robot embodiments and 200 robot control tasks, including 60 atomic manipulation skills such as grasping and placing. Notably, although the dataset provides sufficient diversity, it is not particularly large in scale: the total duration of all manipulation videos is only over 6,000 hours.

4 Experiments

In this section, we introduce comprehensive evaluation experiments, which are conducted to explore the following questions: 1) How does proposed TBot-SA₁ perform compared to existing VLA/WAM models across multiple benchmarks and embodiments? 2) How does the design choice in proposed WSA paradigm contribute to open-world robot learning performance? Next, we first present experimental setup. Then, we demonstrate the merit of TBot-SA₁ through a comparison with state-of-the-arts on both simulation and real-world tasks. Finally, we conduct ablation studies to illustrate major properties of proposed WSA paradigm.

4.1 Experimental Setup

Real-robot setup Real robot experiments are conducted with two embodiments: (1) Unimanual AgileX PiPER with three cameras and 6-DoF robotic arm; (2) Bimanual ARX Lift2 with three cameras and two 7-DoF robotic arms. To ensure diversity of robot behaviors in real-world scenario and investigate cross-embodiment generalization, we design six tabletop manipulation tasks. For each task, objects are randomly initialized within the workspace to introduce variations in object position, orientation, and scene layout. The robot receives language instructions and visual observations as input, and predicts continuous action chunks for closed-loop execution. For testing, we adopt two metrics for comprehensive evaluation of EFMs: (1) Success Rate (SR), the proportion of rollouts in which the task is successfully completed. (2) Completeness (C), the proportion of rollouts in which the sub-task is completed.

Simulation setup We adopt RoboTwin-2.0 benchmark [10] for evaluation in simulation environment. RoboTwin platform is setup with easy and hard modes, where the former offers clean environment without any distracted factor and the latter conduct domain randomization with background changes and distracted objects. Following official protocol [10], 27500 demonstrations (2500 for clean, 25000 for hard) are adopted for model fine-tuning. In addition, we adopt multi-task test setting (i.e., train for once, test for all tasks) to evaluate how much versatility of each model attain. During testing, success rate is adopted as the evaluation metric following the official protocol.

4.2 Main Results on Real-World Tasks

Task description We further evaluate TBot-SA₁ on real-world household manipulation tasks to examine its ability to transfer the learned world-spatial-action representations to physical robot control. The real-world benchmark consists of seven representative tabletop scenarios: RGB blocks ranking, desktop sorting, rubbish cleanup, stationery organization, toys organization, pencil case unzipping and tabletop sweeping. These tasks require the robot to perceive cluttered scenes, interact with objects of different categories and spatial layouts, and complete multi-step manipulation behaviors such as grasping, placing, pushing, and sweeping. Compared with simple pick-and-place tasks, these scenarios place higher demands on spatial reasoning, long-horizon execution, and robustness to diverse object arrangements.

Baselines We compare TBot-SA₁ with three representative baselines: π_0 , a generalist policy for continuous action prediction; $\pi_{0.5}$, a policy for open-world robot control that demonstrates impressive generalization; and InternVLA-A1, a unified model with visual-foresight-guided action generation. For evaluation, we report the final task success rate and the task completeness score. The success rate measures whether the whole task is completed successfully, while the completeness score evaluates partial progress in the task execution.

Evaluation Results As shown in Table 3, TBot-SA₁ achieves a clear lead over baselines in real-world deployment. Across seven household manipulation tasks, TBot-SA₁ achieves an average success rate of 77.38% and an average completeness score of 83.11%, substantially outperforming the $\pi_{0.5}$ baseline, which obtains 55.48% and 59.28%, respectively. These results indicate that the proposed WSA architecture improves both final task completion and intermediate execution quality in real-world manipulation. Moreover, the improvement is consistent across both the single-arm PiPER and bimanual Lift2 platforms, with TBot-SA₁ achieving the highest success rate on every evaluated task.

Table 2: Real-robot experimental setup and task details.




Initial	Finished	Robotic Embodiment	Task Description
		Unimanual: AgileX PiPER (7-DoF)	RGB Block Sorting: Position red block, green block, and blue block from left to right in the specified sequence.
		Unimanual: AgileX PiPER (7-DoF)	Object Organization: Tidy up the table by sorting the toys, stationery, and trash into their designated containers.
		Bimanual: ARX Lift2 (14-DoF)	Trash Cleaning: Pick up the trash with the nearest gripper, hand it over to the other gripper, and place it into the trash bin.
		Bimanual: ARX Lift2 (14-DoF)	Pen Holder Placement: Pick up the markpen and place it into the pen holder.
		Bimanual: ARX Lift2 (14-DoF)	Toy Box Organization: Pick up the toys from the desktop, place them into the toy box, and arrange them neatly inside.
		Bimanual: ARX Lift2 (14-DoF)	Sweep Trash: Sweep the trash into the dustpan using a broom.
		Bimanual: ARX Lift2 (14-DoF)	Unzip Pencil Bag: Grasp the pencil bag, unzip it, and place it back onto the tabletop.

 Table 3: Real-world evaluation results on seven household manipulation tasks. **SR** denotes the task success rate (%), and **C** denotes the task progress score (%).

Method		π_0		$\pi_{0.5}$		InternVLA-A1		TBot-SA ₁	
Type & Size		VLA	3B	VLA	3B	WAM	3B	WSA	3B
<i>Evaluation Results on 7 Real-world Tabletop Robot Manipulation Tasks</i>									
Task Name	Robot Type	SR	C	SR	C	SR	C	SR	C
Rank RGB Blocks	AgileX Piper	70	70	70	70	50	50	80	80
Sort Desktop	AgileX Piper	50	50	77	77	37	37	80	80
Clear Rubbish	ARX Lift2	37	53	77	87	43	58	80	90
Place Stationery	ARX Lift2	7	28	70	77	33	38	87	90
Place Toys	ARX Lift2	13	18	13	27	37	48	60	72
Sweep Trash	ARX Lift2	30	40	27	44	44	78	80	89
Unzip Pencil Case	ARX Lift2	30	41	55	63	25	43	75	75
<i>Average (%)</i>		34	36	55	59	38	53	77	83

These results highlight the effectiveness of the proposed WSA architecture. Rather than relying only on reactive 2D visual-action mapping, TBot-SA₁ couples instruction-aligned visual prediction, 3D world dynamics, and 3D-conditioned action generation, leading to more reliable execution in spatially complex and long-horizon tasks. Notably, this performance is achieved with only moderate-scale pre-training data, suggesting that WSA provides an effective inductive bias for learning transferable real-world manipulation skills under limited demonstration scale.

Table 4: RoboTwin2.0 benchmark experimental results. All models are evaluated under hard setting. **SR** denotes the task success rate (%)

Model Name	Access Status	Model Size	Model Type	SR (%)
Qwen-VLA [33]	Closed-source	5B	VLA	87
Being-H0.7 [24]	Closed-source	3B	WAM	90
MotuBrain [29]	Closed-source	Unknown	WAM	96
π_0 [7]	Open-source	3B	VLA	58
$\pi_{0.5}$ [6]	Open-source	3B	VLA	77
ABot-M0 [38]	Open-source	4B	VLA	80
Motus [4]	Open-source	8B	WAM	87
InternVLA-A1 [9]	Open-source	3B	WAM	88
LingBot-VA [18]	Open-source	7B	WAM	92
Fast-WAM [42]	Open-source	6B	WAM	92
TBot-SA₁(Ours)	Open-source	3B	WSA	93

Table 5: Per-task success rate on RoboTwin2.0 under clean and randomized settings. **SR** denotes the task success rate (%).

Method Type & Size	$\pi_{0.5}$		InternVLA-A1		Motus		FastWAM		TBot-SA ₁	
	VLA	3B	WAM	3B	WAM	8B	WAM	6B	WSA	3B
<i>Evaluation Results on 50 Simulated Bimanual Robot Manipulation Tasks</i>										
Task Name	Clean	Rand.	Clean	Rand.	Clean	Rand.	Clean	Rand.	Clean	Rand.
<i>Place Dual Shoes</i>	75	75	91	88	93	87	94	88	95	94
<i>Move Stapler Pad</i>	56	42	52	67	83	85	77	64	65	75
<i>Stack Blocks Two</i>	97	100	100	100	100	98	100	100	100	100
<i>Scan Object</i>	72	65	86	83	67	66	89	92	85	88
<i>Place Object Stand</i>	91	85	89	88	98	97	90	94	95	95
<i>Place Fan</i>	87	85	98	98	91	87	96	96	98	99
<i>Move Pillbottle Pad</i>	84	61	98	99	93	96	100	99	98	99
<i>Pick Dual Bottles</i>	93	63	92	87	96	90	100	96	94	89
<i>Blocks Ranking Rgb</i>	92	85	92	92	99	97	100	100	95	95
..... (50 tasks)										
<i>Turn Switch</i>	62	54	56	62	84	78	61	59	61	75
<i>Pick Diverse Bottles</i>	81	71	77	77	90	91	80	85	86	76
<i>Place Bread Basket</i>	77	64	90	91	91	94	91	93	99	94
<i>Stack Blocks Three</i>	91	76	90	93	91	95	95	97	98	98
<i>Put Bottles Dustbin</i>	84	79	90	93	81	79	95	90	95	92
<i>Place Can Basket</i>	62	62	81	72	81	76	71	69	80	79
<i>Stamp Seal</i>	79	55	75	75	93	92	90	94	73	82
<i>Hanging Mug</i>	18	17	31	47	38	38	58	62	50	50
<i>Handover Block</i>	66	57	90	73	86	73	95	81	91	84
<i>Stack Bowls Three</i>	77	71	89	90	79	87	80	81	92	93
<i>Place Object Basket</i>	80	76	89	87	81	87	89	88	84	87
<i>Open Microwave</i>	34	77	100	99	95	91	62	45	100	100
<i>Average (%)</i>	83	77	89	88	89	87	92	92	92	93

4.3 Main Results on Simulation Benchmark

RoboTwin2.0 RoboTwin2.0 [25, 10] is a challenging dual-arm manipulation benchmark that evaluates policy generalization under both clean and randomized environments. It covers diverse bimanual skills, including object handover, coordinated manipulation, spatially constrained interaction, and long-horizon tabletop tasks.

As shown in Table 5, TBot-SA₁ achieves strong performance across both settings, obtaining 92% success rate in the clean setting and 93% in the randomized setting. Compared with $\pi_{0.5}$ and Motus, TBot-SA₁ delivers higher average performance, while remaining competitive with FastWAM, a




Figure 4: Investigation of WSA modeling on RoboTwin2.0 Benchmark.




Figure 5: Ablation studies of WSA pre-training on RoboTwin2.0 Benchmark.

representative WAM baseline. More importantly, TBot-SA₁ maintains nearly identical performance between clean and randomized settings, indicating strong robustness to visual and environmental variations. The model also achieves strong results on long-horizon and spatially demanding tasks such as *Move Pillbottle Pad*, *Stack Blocks Three*, and *Open Microwave*. These results suggest that the proposed WSA architecture learns transferable 3D world-action priors that support robust bimanual manipulation under diverse simulation conditions.

4.4 Ablation Study

In this section, we seek to answer the second question via two ablative experiments. First, we perform ablation studies to identify the contribution of causal modeling in WSA, which involves three learning tasks: visual planning, 3D world prediction and 3D action generation. Second, we study the effect of pre-training to reveal the details of full training process in WSA.

Investigation of WSA modeling We adopt the π -style VLA policy proposed in prior works [7, 6] as the baseline model, which integrates the pre-trained QWen3-VL 3B vision-language model [2] with a randomly initialized diffusion transformer serving as the action expert. All ablation experiments are conducted on the RoboTwin2.0 platform, and the corresponding results are presented in Figure 4. Specifically, the baseline model achieves a success rate (SR) of 59.8%. Notably, learning visual planning improves the SR of action generation to 75.7%, while learning 3D world model yields a competitive SR of 74.7%. Full WSA modeling framework further outperforms all individual variants, reaching the best SR of 79.5%. Besides, modeling visual planning is beneficial for manipulation tasks that rely heavily on visual semantics, object appearance, and temporally observable state changes, such as picking bottles and open microwave. In contrast, learning 3D world model improves contact-rich manipulation tasks such as rotate QR Code and stack blocks, which require a deep understanding of spatial geometry, object pose, and hand-object interaction. Based on the above observations, two key conclusions can be drawn. First, learning world dynamics in both 2D and 3D latent space is

essential for generating generalized robot actions. Second, the bidirectional 3D world-action causal mechanism imposes effective constraints on action generation, thereby significantly enhancing the robustness of robot action learning.

The effect of pre-training To comprehensively investigate the efficacy of WSA pre-training, we design three comparative ablation settings: (1) WSA modeling is only applied in the post-training stage without pre-training on diverse heterogeneous robot datasets; (2) WSA modeling is adopted in the pre-training stage, while the post-training process only optimizes the action generation objective; (3) the model is consistently optimized with WSA modeling in both pre-training and post-training stages. As summarized in Figure 5, WSA pre-training delivers a substantial performance gain, increasing the SR from 79.5% to 88.5%. Moreover, applying WSA modeling in both pre-training and post-training further boosts the SR to 91.4%. These results demonstrate that WSA serves as an effective and essential training paradigm for building robust and generalizable embodied foundation models.

4.5 Qualitative Analysis



To reveal how TBot-SA₁ performs robot control, we present qualitative visualizations of the model’s predictions on various tasks. As illustrated in Figure 6, the visualization covers four views: visual observation with text instruction, predicted visual plan, goal 3D world prediction with 3D Gaussian and depth map, and trajectory of executed actions. From the results, we find that predicted visual plans well capture the task objective and well align the generated 3D world. Furthermore, driven by explicit 3D world modeling, the action trajectories are well-aligned with the predicted 3D goal states. These visual results demonstrate that TBot-SA₁ unifies semantic visual planning, causal 3D world dynamics, and closed-loop action generation within a shared latent space. It explicitly infers how actions alter the 3D world and how 3D world states guide action generation, thus leading to robust and generalizable manipulation behaviors.

5 Conclusion

This work addresses critical limitations of current embodied foundation models paradigms, which lack understanding of 3D physical dynamics and bidirectional world–action causality. We propose the World-Spatial-Action (WSA) paradigm and introduce TBot-SA₁, a data-efficient embodied foundation model for generalizable robot control. Built on Mixture-of-Transformers architecture and 3D-centric bidirectional causal mechanism, proposed TBot-SA₁ jointly learns instruction-aligned visual planning, action-conditioned 3D world modeling, and 3D-aware action generation in a unified latent space. The model is pre-trained with only 6,000 hours of demonstration data (1,000 hours from real robots), yet achieves strong performance across simulation and real-world benchmarks. Our work demonstrates that modeling 3D world–action bidirectional causality enables robots to inherently acquire physical commonsense and causal mechanisms, paving a new path for data-efficient and generalizable robot control.

References

- [1] Niket Agarwal, Arslan Ali, Maciej Bala, Yogesh Balaji, Erik Barker, Tiffany Cai, Prithvijit Chattopadhyay, Yongxin Chen, Yin Cui, Yifan Ding, Daniel Dworakowski, Jiaojiao Fan, Michele Fenzi, Francesco Ferroni, Sanja Fidler, Dieter Fox, Songwei Ge, Yunhao Ge, Jinwei Gu, Siddharth Gururani, Ethan He, Jiahui Huang, Jacob Huffman, Pooya Jannaty, Jingyi Jin, Seung Wook Kim, Gergely Klár, Grace Lam, Shiyi Lan, Laura Leal-Taixe, Anqi Li, Zhaoshuo Li, Chen-Hsuan Lin, Tsung-Yi Lin, Huan Ling, Ming-Yu Liu, Xian Liu, Alice Luo, Qianli Ma, Hanzi Mao, Kaichun Mo, Arsalan Mousavian, Seungjun Nah, Sriharsha Niverty, David Page, Despoina Paschalidou, Zeeshan Patel, Lindsey Pavao, Morteza Ramezani, Fitsum Reda, Xiaowei Ren, Vasanth Rao Naik Sabavat, Ed Schmerling, Stella Shi, Bartosz Stefaniak, Shitao Tang, Lyne Tchapti, Przemek Tredak, Wei-Cheng Tseng, Jibin Varghese, Hao Wang, Haoxiang Wang, Heng Wang, Ting-Chun Wang, Fangyin Wei, Xinyue Wei, Jay Zhangjie Wu, Jiashu Xu, Wei Yang, Lin Yen-Chen, Xiaohui Zeng, Yu Zeng, Jing Zhang, Qinsheng Zhang, Yuxuan Zhang, Qingqing Zhao, and Artur Zolkowski. 2025. Cosmos world foundation model platform for physical ai. *arXiv preprint arXiv:2501.03575* (2025).

Current Observation	 Task: "blocks ranking RGB"	 Task: "blocks ranking RGB"
Planned Subgoal Image	  	  
Predicted Goal 3D World	 3D Gaussian  Depth Map	 3D Gaussian  Depth Map
3D Action Execution	  	  



Current Observation	 Task: "stack the bowl on the plate"	 Task: "stack the bowl on the plate"
Planned Subgoal Image	  	  
Predicted Goal 3D World	 3D Gaussian  Depth Map	 3D Gaussian  Depth Map
3D Action Execution	  	  

Figure 6: Visualization examples generated from TBot-SA₁. First row: the current observations with text instruction. Second row: generated subgoal images for visual planning. Third row: predicted 3D world that is rendered via Depth Anything model [20]. Fourth row: the real trajectory of executed actions.

- [2] Shuai Bai, Yuxuan Cai, Ruizhe Chen, Keqin Chen, Xionghui Chen, Zesen Cheng, Lianghao Deng, Wei Ding, Chang Gao, Chunjiang Ge, et al. 2025. Qwen3-vl technical report. *arXiv preprint arXiv:2511.21631* (2025).
- [3] Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibao Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, Humen Zhong, Yuanzhi Zhu, Mingkun Yang, Zhaohai Li, Jianqiang Wan, Pengfei Wang, Wei Ding, Zheren Fu, Yiheng Xu, Jiabo Ye, Xi Zhang, Tianbao Xie, Zesen Cheng, Hang Zhang, Zhibo Yang, Haiyang Xu, and Junyang Lin. 2025. Qwen2.5-VL Technical Report. *arXiv preprint arXiv:2502.13923* (2025).
- [4] Hongzhe Bi, Hengkai Tan, Shenghao Xie, Zeyuan Wang, Shuhe Huang, Haitian Liu, Ruowen Zhao, Yao Feng, Chendong Xiang, Yinze Rong, Hongyan Zhao, Hanyu Liu, Zhizhong Su, Lei Ma, Hang Su, and Jun Zhu. 2025. Motus: A Unified Latent Action World Model. *arXiv preprint arXiv:2512.13030* (2025).

- [5] Johan Bjorck, Fernando Castañeda, Nikita Cherniadev, Xingye Da, Runyu Ding, Linxi "Jim" Fan, Yu Fang, Dieter Fox, Fengyuan Hu, Spencer Huang, Joel Jang, Zhenyu Jiang, Jan Kautz, Kaushil Kundalia, Lawrence Lao, Zhiqi Li, Zongyu Lin, Kevin Lin, Guilin Liu, Edith Llontop, Loic Magne, Ajay Mandlekar, Avnish Narayan, Soroush Nasiriany, Scott Reed, You Liang Tan, Guanzhi Wang, Zu Wang, Jing Wang, Qi Wang, Jiannan Xiang, Yuqi Xie, Yinzheng Xu, Zhenjia Xu, Seonghyeon Ye, Zhiding Yu, Ao Zhang, Hao Zhang, Yizhou Zhao, Ruijie Zheng, and Yuke Zhu. 2026. GR00T N1: An Open Foundation Model for Generalist Humanoid Robots. *arXiv preprint arXiv:2503.14734* (2026).
- [6] Kevin Black, Noah Brown, James Darpinian, Karan Dhabalia, Danny Driess, Adnan Esmail, Michael Robert Equi, Chelsea Finn, Niccolo Fusai, Manuel Y. Galliker, Dibya Ghosh, Lachy Groom, Karol Hausman, brian ichter, Szymon Jakubczak, Tim Jones, Liyiming Ke, Devin LeBlanc, Sergey Levine, Adrian Li-Bell, Mohith Mothukuri, Suraj Nair, Karl Pertsch, Allen Z. Ren, Lucy Xiaoyang Shi, Laura Smith, Jost Tobias Springenberg, Kyle Stachowicz, James Tanner, Quan Vuong, Homer Walke, Anna Walling, Haohuan Wang, Lili Yu, and Ury Zhilinsky. 2025. $\pi_{0.5}$: a Vision-Language-Action Model with Open-World Generalization. In *CoRL*, Vol. 305. 17–40.
- [7] Kevin Black, Noah Brown, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Lachy Groom, Karol Hausman, Brian Ichter, Szymon Jakubczak, Tim Jones, Liyiming Ke, Sergey Levine, Adrian Li-Bell, Mohith Mothukuri, Suraj Nair, Karl Pertsch, Lucy Xiaoyang Shi, James Tanner, Quan Vuong, Anna Walling, Haohuan Wang, and Ury Zhilinsky. 2024. π_0 : A Vision-Language-Action Flow Model for General Robot Control. *arXiv preprint arXiv:2410.24164* (2024).
- [8] Qingwen Bu, Jisong Cai, Li Chen, Xiuqi Cui, Yan Ding, Siyuan Feng, Shenyuan Gao, Xindong He, Xuan Hu, Xu Huang, et al. 2025. Agibot world colosseum: A large-scale manipulation platform for scalable and intelligent embodied systems. *arXiv preprint arXiv:2503.06669* (2025).
- [9] Junhao Cai, Zetao Cai, Jiafei Cao, Yilun Chen, Zeyu He, Lei Jiang, Hang Li, Hengjie Li, Yang Li, Yufei Liu, et al. 2026. InternVLA-A1: Unifying Understanding, Generation and Action for Robotic Manipulation. *arXiv preprint arXiv:2601.02456* (2026).
- [10] Tianxing Chen, Zhanxin Chen, Baijun Chen, Zijian Cai, Yibin Liu, Zixuan Li, Qiwei Liang, Xianliang Lin, Yiheng Ge, Zhenyu Gu, Weiliang Deng, Yubin Guo, Tian Nian, Xuanbing Xie, Qiangyu Chen, Kailun Su, Tianling Xu, Guodong Liu, Mengkang Hu, Huan ang Gao, Kaixuan Wang, Zhixuan Liang, Yusen Qin, Xiaokang Yang, Ping Luo, and Yao Mu. 2025. RoboTwin 2.0: A Scalable Data Generator and Benchmark with Strong Domain Randomization for Robust Bimanual Robotic Manipulation. *arXiv preprint arXiv:2506.18088* (2025).
- [11] Zhe Chen, Jiannan Wu, Wenhui Wang, Weijie Su, Guo Chen, Sen Xing, Muyan Zhong, Qinglong Zhang, Xizhou Zhu, Lewei Lu, et al. 2024. Internvl: Scaling up vision foundation models and aligning for generic visual-linguistic tasks. In *CVPR*. 24185–24198.
- [12] Howard Gardner. 2011. *Frames of Mind: The Theory of Multiple Intelligences*. Basic Books, New York.
- [13] Ryan Hoque, Peide Huang, David J. Yoon, Mouli Sivapurapu, and Jian Zhang. 2025. EgoDex: Learning Dexterous Manipulation from Large-Scale Egocentric Video. *arXiv preprint arXiv:2505.11709* (2025).
- [14] Wenlong Huang, Yu-Wei Chao, Arsalan Mousavian, Ming-Yu Liu, Dieter Fox, Kaichun Mo, and Li Fei-Fei. 2026. PointWorld: Scaling 3D World Models for In-The-Wild Robotic Manipulation. In *CVPR*.
- [15] Moo Jin Kim, Yihuai Gao, Tsung-Yi Lin, Yen-Chen Lin, Yunhao Ge, Grace Lam, Percy Liang, Shuran Song, Ming-Yu Liu, Chelsea Finn, and Jinwei Gu. 2026. Cosmos Policy: Fine-Tuning Video Models for Visuomotor Control and Planning. *arXiv preprint arXiv:2601.16163* (2026).
- [16] Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael Rafailov, Ethan P Foster, Pannag R Sanketi, Quan Vuong, Thomas Kollar, Benjamin Burchfiel, Russ Tedrake, Dorsa Sadigh, Sergey Levine, Percy Liang, and Chelsea Finn. 2025. OpenVLA: An Open-Source Vision-Language-Action Model. In *CoRL*, Vol. 270. 2679–2713.

- [17] Lin Li, Qihang Zhang, Yiming Luo, Shuai Yang, Ruilin Wang, Fei Han, Mingrui Yu, Zelin Gao, Nan Xue, Xing Zhu, Yujun Shen, and Yinghao Xu. 2026. Causal World Modeling for Robot Control. *arXiv preprint arXiv:2601.21998* (2026).
- [18] Lin Li, Qihang Zhang, Yiming Luo, Shuai Yang, Ruilin Wang, Fei Han, Mingrui Yu, Zelin Gao, Nan Xue, Xing Zhu, Yujun Shen, and Yinghao Xu. 2026. Causal World Modeling for Robot Control. *arXiv preprint arXiv:2601.21998* (2026).
- [19] Peiyan Li, Yixiang Chen, Yuan Xu, Jiabing Yang, Xiangnan Wu, Jun Guo, Nan Sun, Long Qian, Xinghang Li, Xin Xiao, Jing Liu, Nianfeng Liu, Tao Kong, Yan Huang, Liang Wang, and Tieniu Tan. 2026. Multi-View Video Diffusion Policy: A 3D Spatio-Temporal-Aware Video Action Model. *arXiv preprint arXiv:2604.03181* (2026).
- [20] Haotong Lin, Sili Chen, Junhao Liew, Donny Y Chen, Zhenyu Li, Guang Shi, Jiashi Feng, and Bingyi Kang. 2025. Depth anything 3: Recovering the visual space from any views. *arXiv preprint arXiv:2511.10647* (2025).
- [21] Yaron Lipman, Ricky TQ Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. 2022. Flow matching for generative modeling. In *ICLR*.
- [22] Songming Liu, Lingxuan Wu, Bangguo Li, Hengkai Tan, Huayu Chen, Zhengyi Wang, Ke Xu, Hang Su, and Jun Zhu. 2025. RDT-1B: a Diffusion Foundation Model for Bimanual Manipulation. In *ICLR*, Vol. 2025. 29982–30009.
- [23] Yunfan Lou, Xiaowei Chi, Xiaojie Zhang, Zezhong Qian, Chengxuan Li, Rongyu Zhang, Yaoxu Lyu, Guoyu Song, Chuyao Fu, Haoxuan Xu, Pengwei Wang, and Shanghang Zhang. 2026. Mask World Model: Predicting What Matters for Robust Robot Policy Learning. *arXiv preprint arXiv:2604.19683* (2026).
- [24] Hao Luo, Wanpeng Zhang, Yicheng Feng, Sipeng Zheng, Haiweng Xu, Chaoyi Xu, Ziheng Xi, Yuhui Fu, and Zongqing Lu. 2026. Being-H0.7: A Latent World-Action Model from Egocentric Videos. *arXiv preprint arXiv:2605.00078* (2026).
- [25] Yao Mu, Tianxing Chen, Zanxin Chen, Shijia Peng, Zhiqian Lan, Zeyu Gao, Zhixuan Liang, Qiaojun Yu, Yude Zou, Mingkun Xu, Lunkai Lin, Zhiqiang Xie, Mingyu Ding, and Ping Luo. 2025. RoboTwin: Dual-Arm Robot Benchmark with Generative Digital Twins. In *CVPR*. 27649–27660.
- [26] Jonas Pai, Liam Achenbach, Victoriano Montesinos, Benedek Forrai, Oier Mees, and Elvis Nava. 2025. mimic-video: Video-Action Models for Generalizable Robot Control Beyond VLAs. *arXiv preprint 2512.15692* (2025).
- [27] Jingjing Qian, Boyao Han, Chen Shi, Lei Xiao, Long Yang, Shaoshuai Shi, and Li Jiang. 2025. GeoPredict: Leveraging Predictive Kinematics and 3D Gaussian Geometry for Precise VLA Manipulation. In *CVPR*.
- [28] Delin Qu, Zeren Gu, Bin Zhao, Dong Wang, and Xuelong Li. 2025. SpatialVLvla:3DVLAA: Exploring Spatial Representations for Visual-Language-Action Models. In *Robotics: Science and Systems (RSS)*.
- [29] MotuBrain Team, Chendong Xiang, Fan Bao, Haitian Liu, Hengkai Tan, Hongzhe Bi, James Li, Jiabao Liu, Jingrui Pang, Kiro Jing, Louis Liu, Mengchen Cai, Rongxu Cui, Ruowen Zhao, Runqing Wang, Shuhe Huang, Yao Feng, Yinze Rong, Zeyuan Wang, and Jun Zhu. 2026. MotuBrain: An Advanced World Action Model for Robot Control. *arXiv preprint arXiv:2604.27792* (2026).
- [30] Octo Model Team, Dibya Ghosh, Homer Walke, Karl Pertsch, Kevin Black, Oier Mees, Sudeep Dasari, Joey Hejna, Tobias Kreiman, Charles Xu, Jianlan Luo, You Liang Tan, Lawrence Yunliang Chen, Pannag Sanketi, Quan Vuong, Ted Xiao, Dorsa Sadigh, Chelsea Finn, and Sergey Levine. 2024. Octo: An Open-Source Generalist Robot Policy. *arXiv preprint arXiv:2405.12213* (2024).

- [31] Team Wan, Ang Wang, Baole Ai, Bin Wen, Chaojie Mao, Chen-Wei Xie, Di Chen, Feiwei Yu, Haiming Zhao, Jianxiao Yang, Jianyuan Zeng, Jiayu Wang, Jingfeng Zhang, Jingren Zhou, Jinkai Wang, Jixuan Chen, Kai Zhu, Kang Zhao, Keyu Yan, Lianghua Huang, Mengyang Feng, Ningyi Zhang, Pandeng Li, Pingyu Wu, Ruihang Chu, Ruili Feng, Shiwei Zhang, Siyang Sun, Tao Fang, Tianxing Wang, et al. 2025. Wan: Open and Advanced Large-Scale Video Generative Models. *arXiv preprint arXiv:2503.20314* (2025).
- [32] Yang Tian, Yuyin Yang, Yiman Xie, Zetao Cai, Xu Shi, Ning Gao, Hangxu Liu, Xuekun Jiang, Zherui Qiu, Feng Yuan, et al. 2025. Interndata-al: Pioneering high-fidelity synthetic data for pre-training generalist policy. *arXiv preprint arXiv:2511.16651* (2025).
- [33] Qiuyue Wang, Mingsheng Li, Jian Guan, Jinhui Ye, Sicheng Xie, Yitao Liu, Junhao Chen, Zhixuan Liang, Jie Zhang, Xintong Hu, Xuhong Huang, Pei Lin, Junyang Lin, Dayiheng Liu, Shuai Bai, Jingren Zhou, Jiazhaio Zhang, Haoqi Yuan, Gengze Zhou, Hang Yin, Ye Wang, Yiyang Huang, Zixing Lei, Wujian Peng, Delin Chen, Yingming Zheng, Jingyang Fan, Xianwei Zhuang, Xin Zhou, Haoyang Li, Anzhe Chen, Tong Zhang, Xuejing Liu, Yuchong Sun, Ruizhe Chen, Zhaohai Li, Chenxu Lü, Zhibo Yang, Tao Yu, and Xionghui Chen. 2026. Qwen-VLA: Unifying Vision-Language-Action Modeling across Tasks, Environments, and Robot Embodiments. *arXiv preprint arXiv:2605.30280* (2026).
- [34] Xuanhan Wang, Huimin Deng, Lianli Gao, and Jingkuan Song. 2025. Scale-Aware Pre-Training for Human-Centric Visual Perception: Enabling Lightweight and Generalizable Models. *arXiv preprint arXiv:2503.08201* (2025).
- [35] Junjie Wen, Yichen Zhu, Jinming Li, Zhibin Tang, Chaomin Shen, and Feifei Feng. 2025. DexVLA: Vision-Language Model with Plug-In Diffusion Expert for General Robot Control. In *CoRL*.
- [36] Adina Yakefu, Bin Xie, Chongyang Xu, Enwen Zhang, Erjin Zhou, Fan Jia, Haitao Yang, Haoqiang Fan, Haowei Zhang, Hongyang Peng, et al. 2025. RoboChallenge: Large-scale Real-robot Evaluation of Embodied Policies. *arXiv preprint arXiv:2510.17950* (2025).
- [37] Ruihan Yang, Qinxu Yu, Yecheng Wu, Rui Yan, Borui Li, An-Chieh Cheng, Xueyan Zou, Yunhao Fang, Xuxin Cheng, Ri-Zhao Qiu, Hongxu Yin, Sifei Liu, Song Han, Yao Lu, and Xiaolong Wang. 2025. EgoVLA: Learning Vision-Language-Action Models from Egocentric Human Videos. *arXiv preprint arXiv:2507.12440* (2025).
- [38] Yandan Yang, Shuang Zeng, Tong Lin, Xinyuan Chang, Dekang Qi, Junjin Xiao, Haoyun Liu, Ronghan Chen, Yuzhi Chen, Dongjie Huo, Feng Xiong, Xing Wei, Zhiheng Ma, and Mu Xu. 2026. ABot-M0: VLA Foundation Model for Robotic Manipulation with Action Manifold Learning. *arXiv preprint arXiv:2602.11236* (2026).
- [39] Seonghyeon Ye, Yunhao Ge, Kaiyuan Zheng, Shenyuan Gao, Sihyun Yu, George Kurian, Suneel Indupuru, You Liang Tan, Chuning Zhu, Jiannan Xiang, Ayaan Malik, Kyungmin Lee, William Liang, Nadun Ranawaka, Jiasheng Gu, Yinzhen Xu, Guanzhi Wang, Fengyuan Hu, Avnish Narayan, Johan Bjorck, Jing Wang, Gwanghyun Kim, Dantong Niu, Ruijie Zheng, Yuqi Xie, Jimmy Wu, Qi Wang, Ryan Julian, Danfei Xu, Yilun Du, Yevgen Chebotar, Scott Reed, Jan Kautz, Yuke Zhu, Linxi "Jim" Fan, and Joel Jang. 2026. World Action Models are Zero-shot Policies. *arXiv preprint arXiv:2602.15922* (2026).
- [40] Zhenhan Yin, Xuanhan Wang, Jiahao Jiang, Kaiyuan Deng, Pengqi Chen, Shuangli Li, Chong Liu, Xing Xu, Jingkuan Song, Lianli Gao, et al. 2025. MiVLA: Towards Generalizable Vision-Language-Action Model with Human-Robot Mutual Imitation Pre-training. *arXiv preprint arXiv:2512.15411* (2025).
- [41] Zhaoshu Yu, Bo Wang, Pengpeng Zeng, Haonan Zhang, Ji Zhang, Zheng Wang, Lianli Gao, Jingkuan Song, Nicu Sebe, and Heng Tao Shen. 2025. A survey on efficient vision-language-action models. *arXiv preprint arXiv:2510.24795* (2025).
- [42] Tianyuan Yuan, Zhibin Dong, Yicheng Liu, and Hang Zhao. 2026. Fast-WAM: Do World Action Models Need Test-time Future Imagination? *arXiv preprint arXiv:2603.16666* (2026).

- [43] Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. 2023. Sigmoid Loss for Language Image Pre-Training. In *ICCV*. 11975–11986.
- [44] Peng-Fei Zhang, Ying Cheng, Xiaofan Sun, Shijie Wang, Fengling Li, Lei Zhu, and Heng Tao Shen. 2025. A step toward world models: A survey on robotic manipulation. *arXiv preprint arXiv:2511.02097* (2025).
- [45] Wenyao Zhang, Hongsi Liu, Zekun Qi, Yunnan Wang, Xinqiang Yu, Jiazhao Zhang, Runpei Dong, Jiawei He, Fan Lu, He Wang, Zhizheng Zhang, Li Yi, Wenjun Zeng, and Xin Jin. 2025. DreamVLA: A Vision-Language-Action Model Dreamed with Comprehensive World Knowledge. In *NeurIPS*.
- [46] Wenyao Zhang, Hongsi Liu, Zekun Qi, Yunnan Wang, Xinqiang Yu, Jiazhao Zhang, Runpei Dong, Jiawei He, He Wang, Zhizheng Zhang, Li Yi, Wenjun Zeng, and Xin Jin. 2025. DreamVLA: Dream Comprehensive World Knowledge for Vision-Language-Action Model. In *NeurIPS*.
- [47] Qingqing Zhao, Yao Lu, Moo Jin Kim, Zipeng Fu, Zhuoyang Zhang, Yecheng Wu, Zhaoshuo Li, Qianli Ma, Song Han, Chelsea Finn, Ankur Handa, Tsung-Yi Lin, Gordon Wetzstein, Ming-Yu Liu, and Donglai Xiang. 2025. CoT-VLA: Visual Chain-of-Thought Reasoning for Vision-Language-Action Models. In *CVPR*. 1702–1713.
- [48] Haoyu Zhen, Xiaowen Qiu, Peihao Chen, Jincheng Yang, Xin Yan, Yilun Du, Yining Hong, and Chuang Gan. 2024. 3D-VLA: A 3D Vision-Language-Action Generative World Model. 61229–61245.
- [49] Brianna Zitkovich, Tianhe Yu, Sichun Xu, Peng Xu, Ted Xiao, Fei Xia, Jialin Wu, Paul Wohlhart, Stefan Welker, Ayzaan Wahid, Quan Vuong, Vincent Vanhoucke, Huong Tran, Radu Soricut, Anikait Singh, Jaspiar Singh, Pierre Sermanet, Pannag R. Sanketi, Grecia Salazar, Michael S. Ryoo, Krista Reymann, Kanishka Rao, Karl Pertsch, Igor Mordatch, Henryk Michalewski, Yao Lu, Sergey Levine, Lisa Lee, Tsang-Wei Edward Lee, Isabel Leal, Yuheng Kuang, Dmitry Kalashnikov, Ryan Julian, Nikhil J. Joshi, Alex Irpan, Brian Ichter, Jasmine Hsu, Alexander Herzog, Karol Hausman, Keerthana Gopalakrishnan, Chuyuan Fu, Pete Florence, Chelsea Finn, Kumar Avinava Dubey, Danny Driess, Tianli Ding, Krzysztof Marcin Choromanski, Xi Chen, Yevgen Chebotar, Justice Carbajal, Noah Brown, Anthony Brohan, Montserrat Gonzalez Arenas, and Kehang Han. 2023. RT-2: Vision-Language-Action Models Transfer Web Knowledge to Robotic Control. In *CoRL*, Vol. 229. 2165–2183.

Task Name	Clean	Random
adjust_bottle	100	99
beat_block_hammer	98	95
blocks_ranking_rgb	95	95
blocks_ranking_size	92	92
click_alarmclock	100	100
click_bell	100	100
dump_bin_bigbin	97	99
grab_roller	100	100
handover_block	91	84
handover_mic	100	100
hanging_mug	50	50
lift_pot	100	99
move_can_pot	86	95
move_pillbottle_pad	99	100
move_playingcard_away	100	100
move_stapler_pad	65	75
open_laptop	100	100
open_microwave	100	100
pick_diverse_bottles	86	76
pick_dual_bottles	94	89
place_a2b_left	92	97
place_a2b_right	94	94
place_bread_basket	99	94
place_bread_skillet	90	89
place_burger_fries	100	100
place_can_basket	80	79
place_cans_plasticbox	100	99
place_container_plate	99	100
place_dual_shoes	95	94
place_empty_cup	100	99
place_fan	98	99
place_mouse_pad	86	86
place_object_basket	84	87
place_object_scale	85	88
place_object_stand	95	95
place_phone_stand	99	97
place_shoe	99	99
press_stapler	83	78
put_bottles_dustbin	95	92
put_object_cabinet	90	93
rotate_qrcode	87	92
scan_object	85	88
shake_bottle	100	100
shake_bottle_horizontally	100	100
stack_blocks_three	98	98
stack_blocks_two	100	100
stack_bowls_three	92	93
stack_bowls_two	97	100
stamp_seal	73	82
turn_switch	61	75
AVG.	92	93

Table 6: Per-task success rates on RoboTwin2.0 under clean and randomized evaluation settings.